



# On the Origin of the Programming-Models

Tim Mattson

Intel, Parallel Computing Lab



# Legal Disclaimer



INFORMATION IN THIS DOCUMENT IS PROVIDED IN CONNECTION WITH INTEL PRODUCTS. NO LICENSE, EXPRESS OR IMPLIED, BY ESTOPPEL OR OTHERWISE, TO ANY INTELLECTUAL PROPERTY RIGHTS IS GRANTED BY THIS DOCUMENT. EXCEPT AS PROVIDED IN INTEL'S TERMS AND CONDITIONS OF SALE FOR SUCH PRODUCTS, INTEL ASSUMES NO LIABILITY WHATSOEVER AND INTEL DISCLAIMS ANY EXPRESS OR IMPLIED WARRANTY, RELATING TO SALE AND/OR USE OF INTEL PRODUCTS INCLUDING LIABILITY OR WARRANTIES RELATING TO FITNESS FOR A PARTICULAR PURPOSE, MERCHANTABILITY, OR INFRINGEMENT OF ANY PATENT, COPYRIGHT OR OTHER INTELLECTUAL PROPERTY RIGHT.

A "Mission Critical Application" is any application in which failure of the Intel Product could result, directly or indirectly, in personal injury or death. SHOULD YOU PURCHASE OR USE INTEL'S PRODUCTS FOR ANY SUCH MISSION CRITICAL APPLICATION, YOU SHALL INDEMNIFY AND HOLD INTEL AND ITS SUBSIDIARIES, SUBCONTRACTORS AND AFFILIATES, AND THE DIRECTORS, OFFICERS, AND EMPLOYEES OF EACH, HARMLESS AGAINST ALL CLAIMS COSTS, DAMAGES, AND EXPENSES AND REASONABLE ATTORNEYS' FEES ARISING OUT OF, DIRECTLY OR INDIRECTLY, ANY CLAIM OF PRODUCT LIABILITY, PERSONAL INJURY, OR DEATH ARISING IN ANY WAY OUT OF SUCH MISSION CRITICAL APPLICATION, WHETHER OR NOT INTEL OR ITS SUBCONTRACTOR WAS NEGLIGENT IN THE DESIGN, MANUFACTURE, OR WARNING OF THE INTEL PRODUCT OR ANY OF ITS PARTS.

Intel may make changes to specifications and product descriptions at any time, without notice. Designers must not rely on the absence or characteristics of any features or instructions marked "reserved" or "undefined". Intel reserves these for future definition and shall have no responsibility whatsoever for conflicts or incompatibilities arising from future changes to them. The information here is subject to change without notice. Do not finalize a design with this information. The products described in this document may contain design defects or errors known as errata which may cause the product to deviate from published specifications. Current characterized errata are available on request.

Contact your local Intel sales office or your distributor to obtain the latest specifications and before placing your product order.

Copies of documents which have an order number and are referenced in this document, or other Intel literature, may be obtained by calling 1-800-548-4725, or go to: <http://www.intel.com/design/literature.htm>

Knights Landing and other code names featured are used internally within Intel to identify products that are in development and not yet publicly announced for release. Customers, licensees and other third parties are not authorized by Intel to use code names in advertising, promotion or marketing of any product or services and any such use of Intel's internal code names is at the sole risk of the user

Intel, Cilk, VTune, Xeon, Xeon Phi and the Intel logo are trademarks of Intel Corporation in the United States and other countries.

\*Other names and brands may be claimed as the property of others.

Copyright ©2013 Intel Corporation.

## Optimization Notice

Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice.

Notice revision #20110804

# Human Readable Disclaimer



- The views expressed in this talk are those of the speaker and not his employer.
- These slides are borrowed from other talks and other people. Sorry if you've heard much of this before. But remember, I'm talking as an HPC historian in this talk, and we know that ...
  - While history repeats itself, historians often repeat each other.

I work in Intel's research labs. I don't build products. Instead, I get to poke into dark corners and think silly thoughts... just to make sure we don't miss any great ideas.

Hence, my views are by design far "off the roadmap".

# The quest for the “right” Programming model

- I am a molecular physicist .... I don't even like computers ... So its ironic that I am best known for my work on programming models:

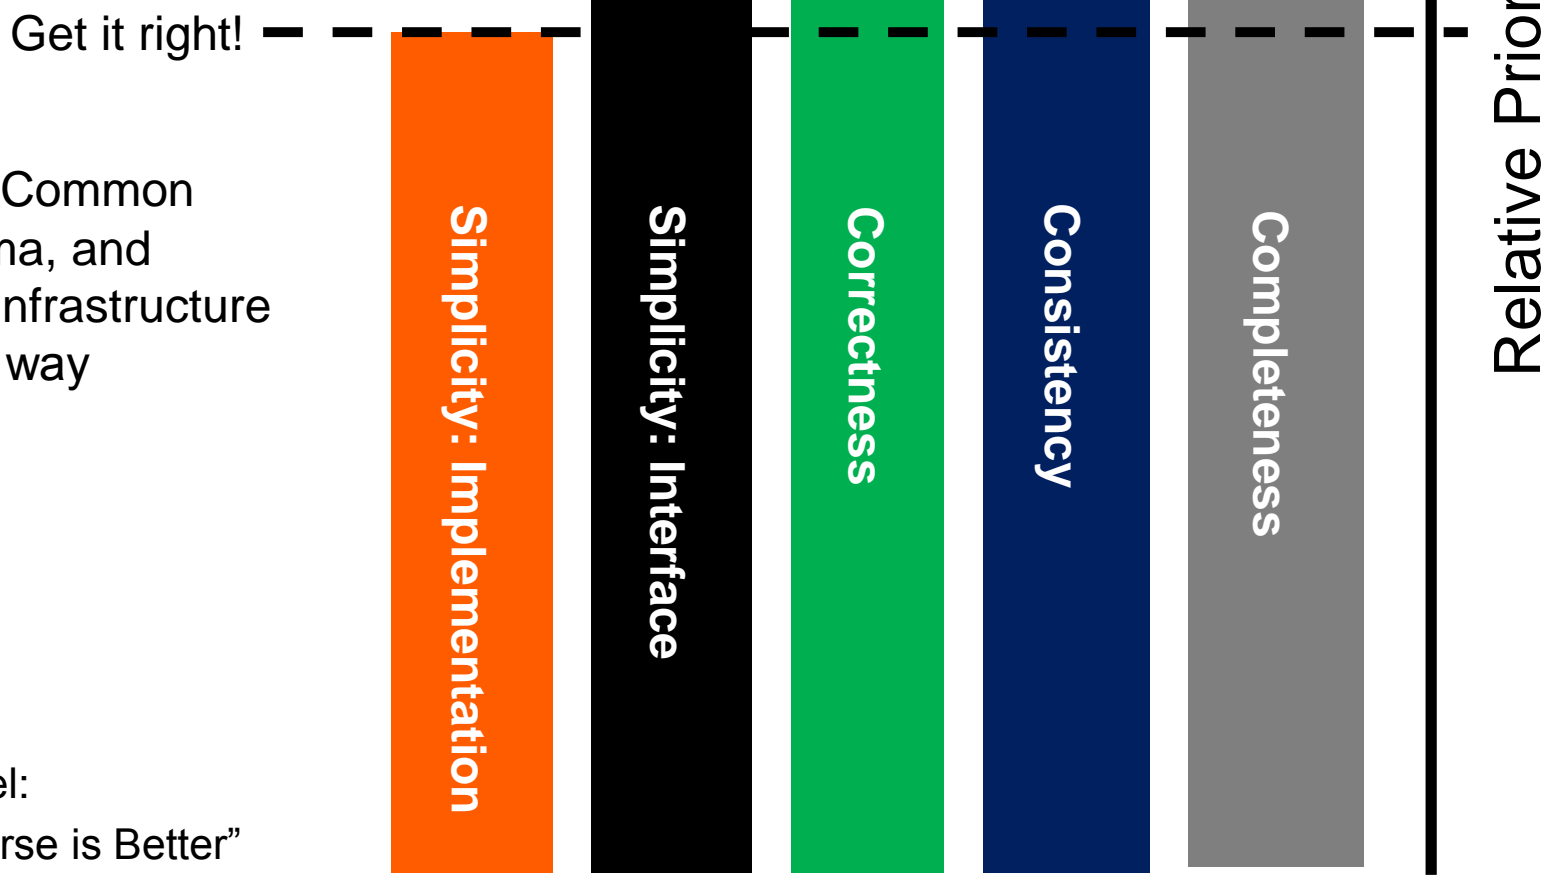
Ada	mid-80's	DOD mandated object based language	Complex, big, and nobody likes to be forced to use a new language
Strand	Late 1980's	Concurrent logic programming	Failed miserably ... nobody wanted to learn a new language
Linda	Late 1980's to mid-90's	Coordination language built around a tuplespace	Failed miserably ... why pay money when you can get something (PVM) that works for free?
MPI	mid-90's to current	Roll the best message passing ideas into one system	MPI is the most successful parallel programming language ever.
OpenMP	Late-90's to present	Shared memory programming made simple	The second most successful parallel programming language.
OpenCL	2007 to present	Portable platform for heterogeneous systems	Application programmers screwed up ... long live CUDA
OCR	2010 to present	Assync. Multi-tasking.	Too early to say ....

# Making sense of the programming models

- To understand which programming models succeed and which fail, we need to start with the famous essay by Richard Gabriel ... “The rise of worse is better”
  - An essay that tried to explain the failure of common LISP to become a dominant programming model.

# Design Philosophy: “The Right Thing”

**Example:** Common  
Lisp, Schema, and  
supporting infrastructure  
... The MIT way



Richard Gabriel:  
The rise of Worse is Better”

“<https://www.jwz.org/doc/worse-is-better.html>

Third party names are the property of their owners

# Design Philosophy: “Worse is Better”

Get it right!

**Example:** Unix and C  
... The New Jersey way

Simplicity: Implementation

Simplicity: Interface

Correctness

Consistency

Completeness

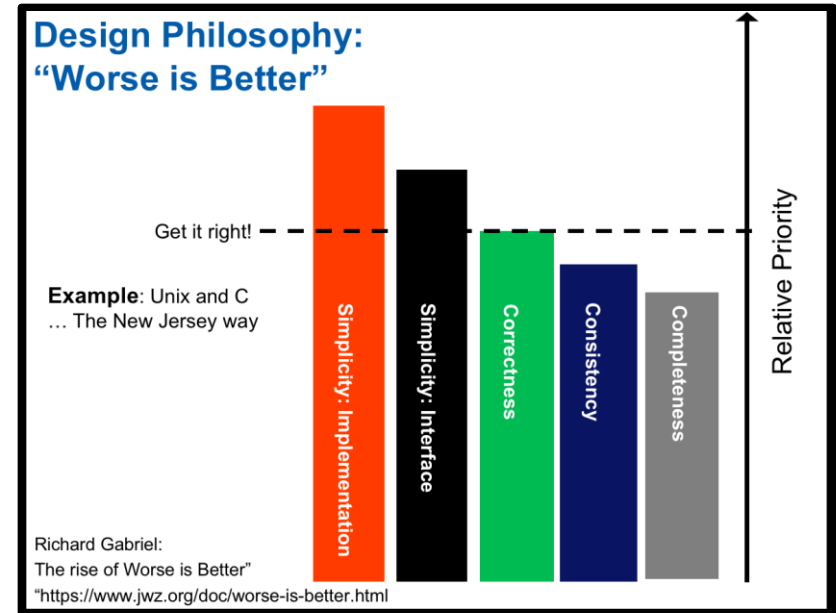
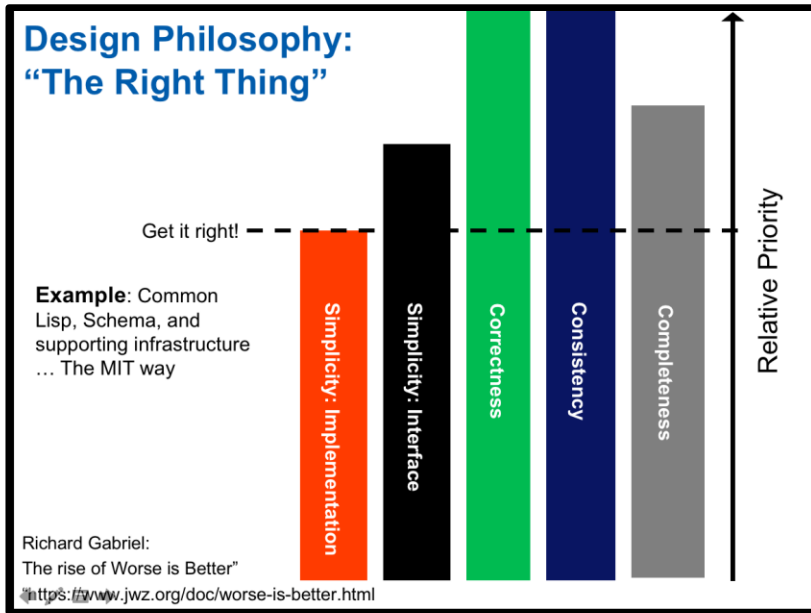
Relative Priority

Richard Gabriel:  
The rise of Worse is Better”

“<https://www.jwz.org/doc/worse-is-better.html>”

Third party names are the property of their owners

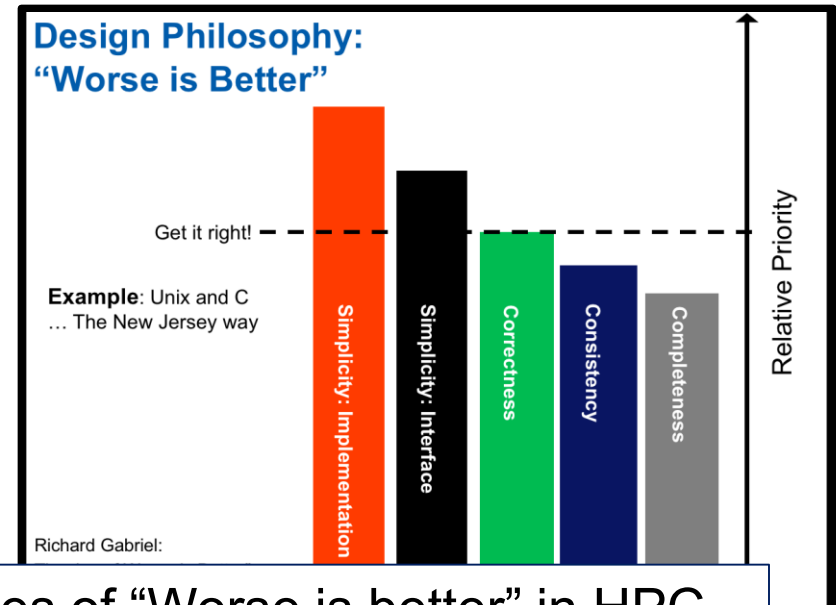
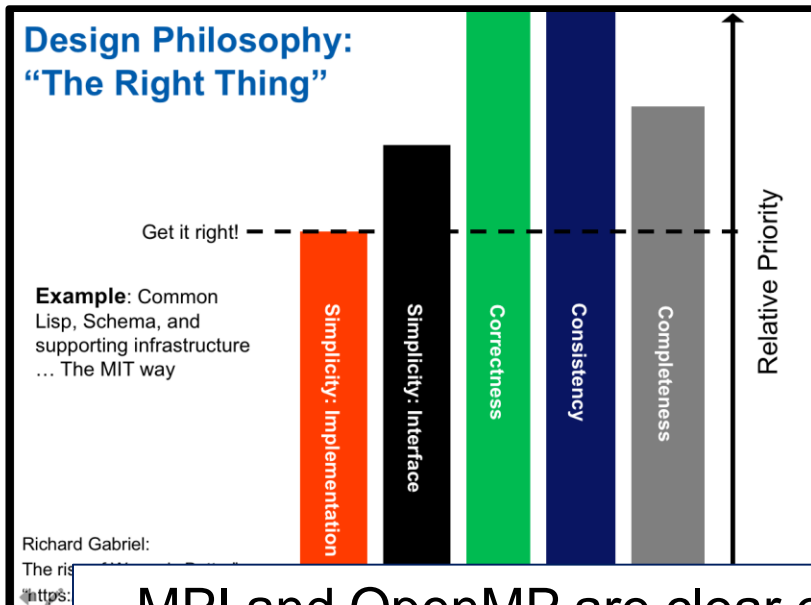
# Which Design Philosophy wins?



- History shows again and again ... "Worse is better".
  - While "the right thing" take the time to "get it right", the "worse is better" folks are busy establishing a user base.
  - "Worse is better" programmers are conditioned to sacrifice safety, convenience, and hassle to get good performance.
  - Since "worse is better" stresses implementation simplicity, its available everywhere.
  - With a large user base, once "worse is better" has spread, there is pressure to improve it ... so over time it becomes good enough



# Which Design Philosophy wins?



MPI and OpenMP are clear examples of "Worse is better" in HPC

- History shows again and again ... "Worse is better".
  - While "the right thing" take the time to "get it right", the "worse is better" folks are busy establishing a user base.
  - "Worse is better" programmers are conditioned to sacrifice safety, convenience, and hassle to get good performance.
  - Since "worse is better" stresses implementation simplicity, its available everywhere.
  - With a large user base, once "worse is better" has spread, there is pressure to improve it ... so over time it becomes good enough

# But there must be better ways to program parallel computers!

## Parallel programming environments in the 90's

ABCPL	CORRELATE	GLU	Mentat	Parafraze2	pC++
ACE	CPS	GUARD	Legion	Paralation	SCHEDULE
ACT++	CRL	HASL	Meta Chaos	Parallel-C++	SciTL
Active messages	CSP	Haskell	Midway	Parallaxis	POET
Adl	Cthreads	HPC++	Millipede	ParC	SDDA
Adsmith	CUMULVS	JAVAR.	CparPar	ParLib++	SHMEM
ADDAP	DAGGER	HORUS	Mirage	ParLin	SIMPLE
AFAPI	DAPPLE	HPC	MpC	Parmacs	Sina
ALWAN	Data Parallel C	IMPACT	MOSIX	Parti	SISAL
AM	DC++	ISIS	Modula-P	pC	distributed smalltalk
AMDC	DCE++	JAVAR	Modula-2*	pC++	SMI
AppLeS	DDD	JADE	Multipol	PCN	SONiC
Amoeba	DICE	Java RMI	MPI	PCP:	Split-C.
ARTS	DIPC	javaPG	MPC++	PH	SR
Athapascan-0b	DOLIB	JavaSpace	Munin	PEACE	Sthreads
Aurora	DOME	JIDL	Nano-Threads	PCU	Strand.
Automap	DOSMOS.	Joyce	NESL	PET	SUIF.
bb_threads	DRL	Khoros	NetClasses++	PETSc	Synergy
Blaze	DSM-Threads	Karma	Nexus	PENNY	Telegrphos
BSP	Ease .	KOAN/Fortran-S	Nimrod	Phosphorus	SuperPascal
BlockComm	ECO	LAM	NOW	POET.	TCGMSG.
C*.	Eiffel	Lilac	Objective Linda	Polaris	Threads.h++.
"C* in C	Eilean	Linda	Occam	POOMA	TreadMarks
C**	Emerald	JADA	Omega	POOL-T	TRAPPER
CarlOS	EPL	WWWinda	OpenMP	PRESTO	uC++
Cashmere	Excalibur	ISETL-Linda	Orca	P-RIO	UNITY
C4	Express	ParLin	OOF90	Prospero	UC
CC++	Falcon	Eilean	P++	Proteus	V
Chu	Filaments	P4-Linda	P3L	QPC++	ViC*
Charlotte	FM	Glenda	p4-Linda	PVM	Visifold V-NUS
Charm	FLASH	POSYBL	Pablo	PSI	VPE
Charm++	The FORCE	Objective-Linda	PADE	PSDM	Win32 threads
Cid	Fork	LiPS	PADRE	Quake	WinPar
Cilk	Fortran-M	Locust	Panda	Quark	WWWinda
CM-Fortran	FX	Lparx	Papers	Quick Threads	XENOOKS
Converse	GA	Lucid	AFAPI.	Sage++	XPC
Code	GAMMA	Maisie	Para++	SCANDAL	Zounds
COOL	Glenda	Manifold	Paradigm	SAM	ZPL

Third party names are the property of their owners.

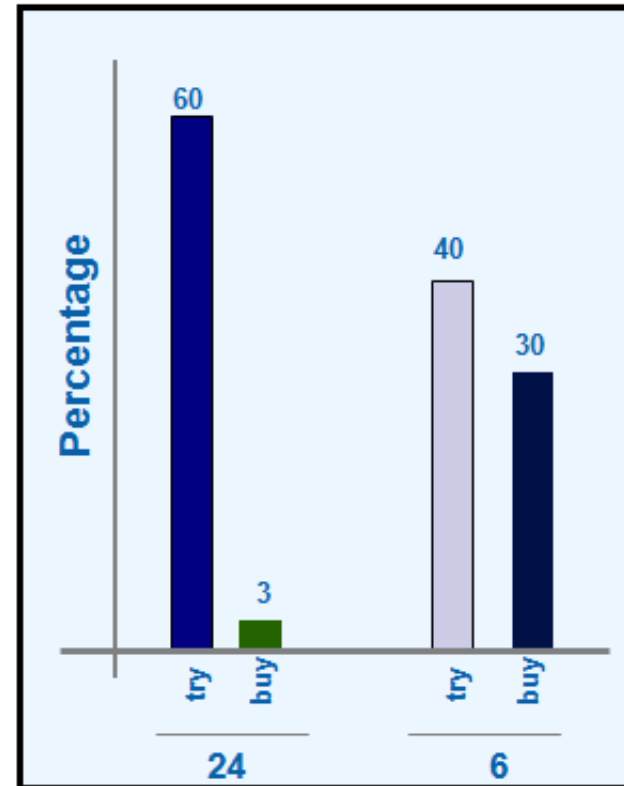
# A warning I've been making for the last 10 years

Is it bad to have so many languages?

Too many options can hurt you

## ■ The Draeger Grocery Store experiment consumer choice:

- Two Jam-displays with coupon's for purchase discount.
  - 24 different Jam's
  - 6 different Jam's
- How many stopped by to try samples at the display?
- Of those who "tried", how many bought jam?



Programmers don't need a glut of options ... just give us something that works OK on every platform we care about. Give us a decent standard and we'll do the rest

The findings from this study show that an extensive array of options can at first seem highly appealing to consumers, yet can reduce their subsequent motivation to purchase the product.

Iyengar, Sheena S., & Lepper, Mark (2000). When choice is demotivating: Can one desire too much of a good thing? *Journal of Personality and Social Psychology*, 76, 995-1006.

# But this isn't quite fair

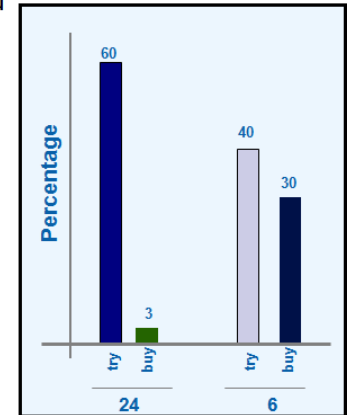
- The glut of parallel languages and choice overload matter a great deal when your job is to reach out to application programmers and software vendors.
- But is it really something Computer Science researchers should be overly concerned about?

We tried to solve the programmability problem by searching for the right programming environment  
Parallel programming environments in the 90's

ABCPL	CORRELATE	GLU	Mentat	Parafraze2	pC++
ACE	CPH	GUARD	Legion	Paralation	SCHEDULE
ACT++	CRL	HALE	Meta Chaos	Parallel C++	SoTL
Active messages	CSP	Hasdel	Midway	Parallaxis	SoTL
AEL	Cthreads	HPC++	Millipede	ParC	POST
Adsmith	CUMULVS	JAVAR	OpazPar	ParLib++	RDMA
ADDAP	DAGGER	ROBUST	Mirage	ParLin	SHMEM
AFAPI	DAPPLE	HPC	Mpc	Parmacs	SIMPLE
ALWAN	Data Parallel C	DMPACT	MODX	Parli	Sisal
AM	DC++	ISIS	Modula-2*	pC	distributed smalltalk
AMDC	DDD	JAVAR	Modula-2*	pC++	SML
AppLe8	DADE	JADE	Modula-2*	PCN	SONIC
Armoeba	DICE	Java RMI	MPI	PCP	Split-C
ARTS	DIPC	javaPG	MPC++	PH	SR
Athapascan-0b	DOLIB	JavaSpace	Monin	PEACE	Strand
Aurora	DOMS	JIDL	Nano-Threads	PCU	SUP
Automap	DOSEOS	Joyce	NESL	PET	Synegy
M2-threads	DRL	Khoros	NarClass++	PETs	Telegraph
Blaze	DSM-Threads	Karma	Nexus	PENNY	SuperPascal
BSP	Ease	KOAN/Fortran-3	Nimrod	Phosphorus	
BlockComm					
C*					
C* in C					
C++					
CarOS					
Cashmere					
C4					
CC++					
Chi					
Charlotte					
Charm					
Charm++					
Cid					
Clic					
CM-Fortran					
Comverse					
Code					
COOL					

Is it bad to have so many languages?  
Too many options can hurt you

- The Draeger Grocery Store experiment consumer choice:
  - Two Jam-displays with coupon's for purchase discount.
    - 24 different Jam's
    - 6 different Jam's
  - How many stopped by to try samples at the display?
  - Of those who "tried", how many bought jam?



Programmers don't need a glut of options ... just give us something that works OK on every platform we care about. Give us a decent standard and we'll do the rest

The findings from this study show that an extensive array of options can at first seem highly appealing to consumers, yet can reduce their subsequent motivation to purchase the product.  
Iyengar, Sheena S., & Lepper, Mark (2000). When choice is demotivating: Can one desire too much of a good thing? *Journal of Personality and Social Psychology*, 76, 995-1006.

We need a more nuanced way to think about the role of academic computer science research

# Evolution

- Two Major Models of evolution:

- **Phyletic Gradualism:**

A slow, continuous and gradual process of change.



How these trends would appear in the fossil record



time

- **Punctuated Equilibrium:**

Long periods where little change is observed interspersed with periods of abrupt change.



time

Source: <http://en.wikipedia.org/wiki/File:Punctuated-equilibrium.svg>

What actually happens in biological Evolution?

# What actually happens in Biology?

- Both Phyletic Gradualism and Punctuated Equilibrium occur:
  - Decompose the ecosystem into relatively isolated niches.
  - Inside an isolated niche, phyletic Gradualism occurs
  - An event changes the ecosystem (e.g. climate change) so what was once an obscure Niche becomes the new mainstream.
  - Since an adaptive species evolved to match the “new mainstream” in a protected niche, it is now well positioned to dominate the new normal.
- So the answer of phyletic gradualism vs. punctuated equilibrium is “both”.

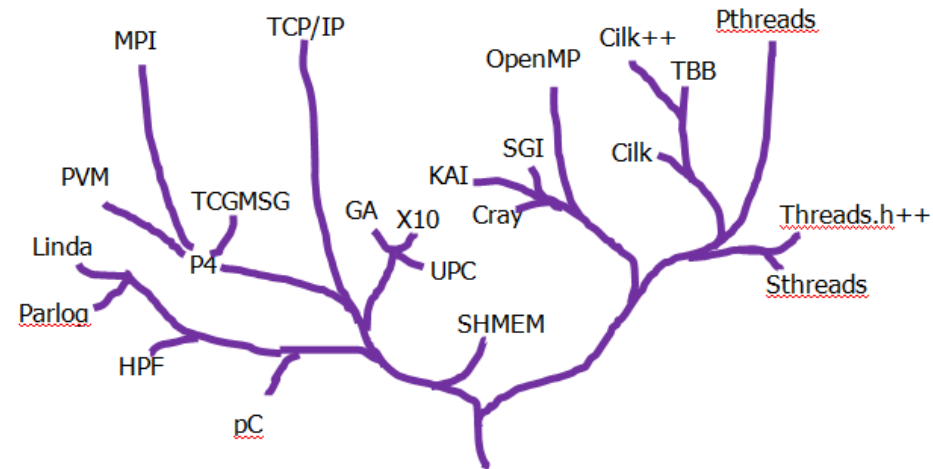
How does this apply to programming models?

# Evolution of parallel programming environments

- Both Models are observed depending on the ecosystem:

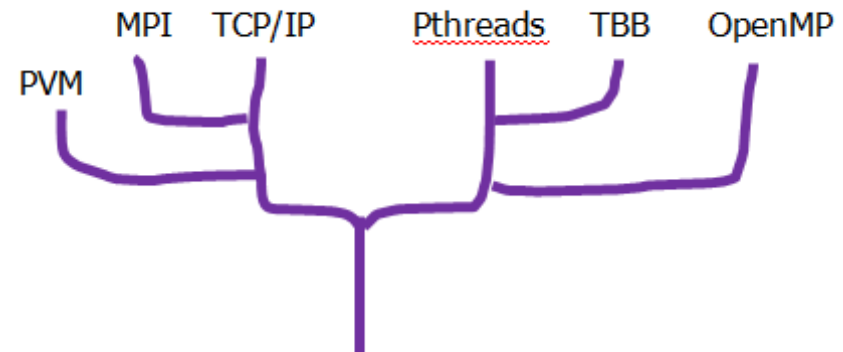
- Research community**

- Phyletic Gradualism with lateral gene transfer:
  - (1) large, growing population of parallel programming systems,
  - (2) occupy isolated environmental niches, and
  - (3) few real users.



- Professional Application developers**

- Punctuated equilibrium ...
  - (1) a small number of stable programming systems that
  - (2) change rarely in response to abrupt external forces.



Computer Science research creates the innovation so new "species" are ready when environmental conditions change and new "mainstream" systems emerge.

# External change and the evolution

- External changes establish a “new normal” and fringe academic projects walk in to dominate.
- So It’s a “good thing” that researchers are creating so many programming languages ... as long as they don’t prematurely push them into the mainstream (choice overload).
- To understand the future of computing we need to look at the key changes (or inflection points) just around the corner
- And then ask ... which academic/research projects are “in the wings” waiting to emerge and dominate the new normal

We tried to solve the programmability problem by searching for the right programming environment

Parallel programming environments in the 90’s

ABCPL	CORRELATE	GLU	Mentat	Parafraze2	pC++
ACE	CPS	GUARD	Legion	Paralation	SCHEDULE
ACT++	CRL	HAUL	Meta Chaos	Parallel-C++	SCITL
Active messages	CSP	Haskell	Midway	Parallaxis	POET
Adl	Cthreads	HPC++	Millipede	ParC	SDDA
Ademith	CTMULV3	JAVAR	CpuPar	ParLb++	SHMEM
ADDP	DAGGER	HORUS	Mirage	ParLin	SIMPLE
AFAPI	DAPPLE	HPC	MpC	Parmaes	Sino
ALWAN	Data Parallel C	IMPACT	MOBIX	Parti	SISAL
AM	DC++	ISB	Module-P	pC	distributed smalltalk
AMDC	DCE++	JAVAR	Module-2*	pC++	SML
AppLeS	DDD	JADE	Multipol	PCN	SOIUC
Amosha	DICE	Java RMI	NPI	PCP	Split-C
ARTS	DIPC	javaPG	MPC++	PH	SR
Athapascan-0b	DOLIB	JavaSpace	Manin	PEACE	Sthreads
Aurora	DOMIE	JIDL	Nano-Threads	PCU	Strand
Automap	DOIMOS	Joyce	NEL	PET	SUF
W_threads	DRL	Khoros	NetClasses++	PETSc	Synergy
Blaize	DSM-Threads	Karna	Nexus	PENNY	Telegraphos
BBP	Ease	KOANFortran-S	Nimrod	Phosphorus	SuperPascal
BlockComm	ECO	LAM	NOW	POET	TCOMEG
C*	Effiel	Lilac	Objective Linda	Polaris	Threads.h++
C** in C	Eilean	Linda	Occam	POOMA	TreadMarks
C**	Emerald	JADA	Omega	POOL-T	TRAPPER
CarLOS	EPL	WWWinda	OpenMP	PRESTO	UC++
Cashmere	Excalibur	IBETL-Linda	Oca	P-RIO	UNITY
C4	Express	ParLin	OOP90	Prospero	UC
CC++	Falcon	Eilean	P++	Proteus	V
Chu	Filaments	P4-Linda	P3L	QPC++	Vic*
Charlotte	FM	Glenda	p4-Linda	PVM	Vinifold V-NUS
Charm	FLASH	Pablo	PPI	QPC++	VPE
Charm++	TheFORCE	Objective-Linda	PADRE	QPC++	Win32 threads
Cid	Fork	LIPS	PADRE	QPC++	WinPar
Clik	Fortran-M	Locust	Panda	Quark	WWWinda
CN-Fortran	FX	Lparx	Papers	Quick Threads	XENOOOS
Converse	GA	Liquid	AFAPI	Sage++	XPC
Code	GAMDA	Maiaie	Para++	SCANDAL	Zonda
COOL	Glenda	Manifold	Paradigm	SAM	ZPL

Third party names are the property of their owners.

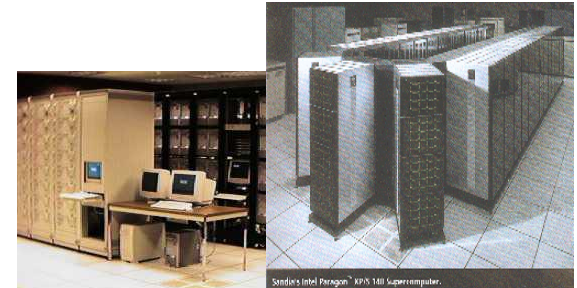


# Key HPC hardware inflection points

- The first multiprocessor: Burroughs B5000, 1961
- SMP goes mainstream: the Intel® Pentium™ technology in 1995 (up to two processors) and the Intel® Pentium Pro® processors (up to four).
- MPPs (e.g. Intel ® Paragon, TMC CM5, Cray T3D) in early 90's,
- Clusters (Stacked Sparc pizza boxes late 80's) and Linux clusters starting with Beowulf in 1994.
- GPGPU programming starts in early 2000's but using primitive shader language
- NVIDIA innovations lead to fully programmable GPUs



Dual socket Pentium pro board (~1997)



NCSA super-cluster (1998) and Paragon XPS 140 (1994)



NVIDIA GeForce 8800/HD2900 (~2006)

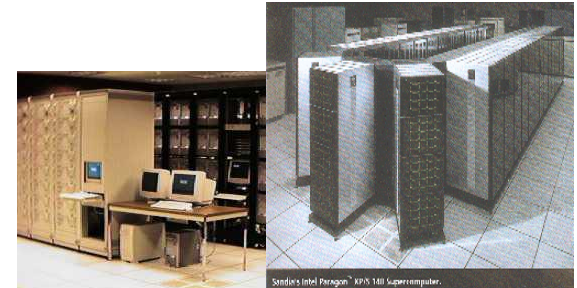
# Key HPC hardware inflection points

- The first multiprocessor technology in 1961
- SMP goes mainstream technology in 1995 (up to two processors) Pro (up to four processors).



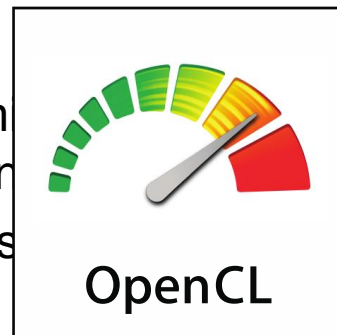
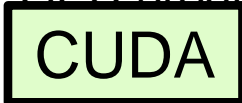
Dual socket Pentium pro board (~1997)

- MPPs (e.g. Paragon TMC GM5, Cray T3D) in early 90's,
- Clusters (Stacked Sparc pizza boxes late 80's) and Linux clusters starting with Beowulf in 1994.



NCSA super-cluster (1998) and Paragon XPS 140 (1994)

- GPGPU programming primarily in the early 2000's but using
- NVIDIA innovations in programmable GPUs



NVIDIA GeForce 8800/HD2900 (~2006)

# Adoption of HPC programming models

- New HPC programming models only occur at hardware inflection points
- In between inflection points, it is more productive to improve existing models than to push new ones.

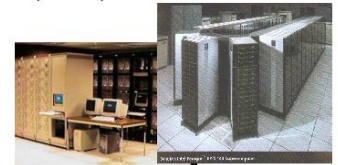
## Key HPC hardware inflection points

- The first multiprocessor technology in 1961
- SMP goes mainstream in 1995 (up to two processors) and OpenMP Pro (up to four processors).



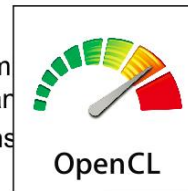
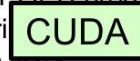
Dual socket Pentium pro board (~1997)

- MPPs (e.g. Paragon T3E, Cray T3D) in early 90's,
- Clusters (Stacked Sparc pizza boxes late 80's) and Linux clusters starting with Beowulf in 1994.



NCSA super-cluster (1998) and Paragon XPS 140 (1994)

- GPGPU programming primarily in the early 2000's but using CUDA
- NVIDIA innovations in programmable GPUs



NVIDIA GeForce 8800/HD2900 (~2006)

16

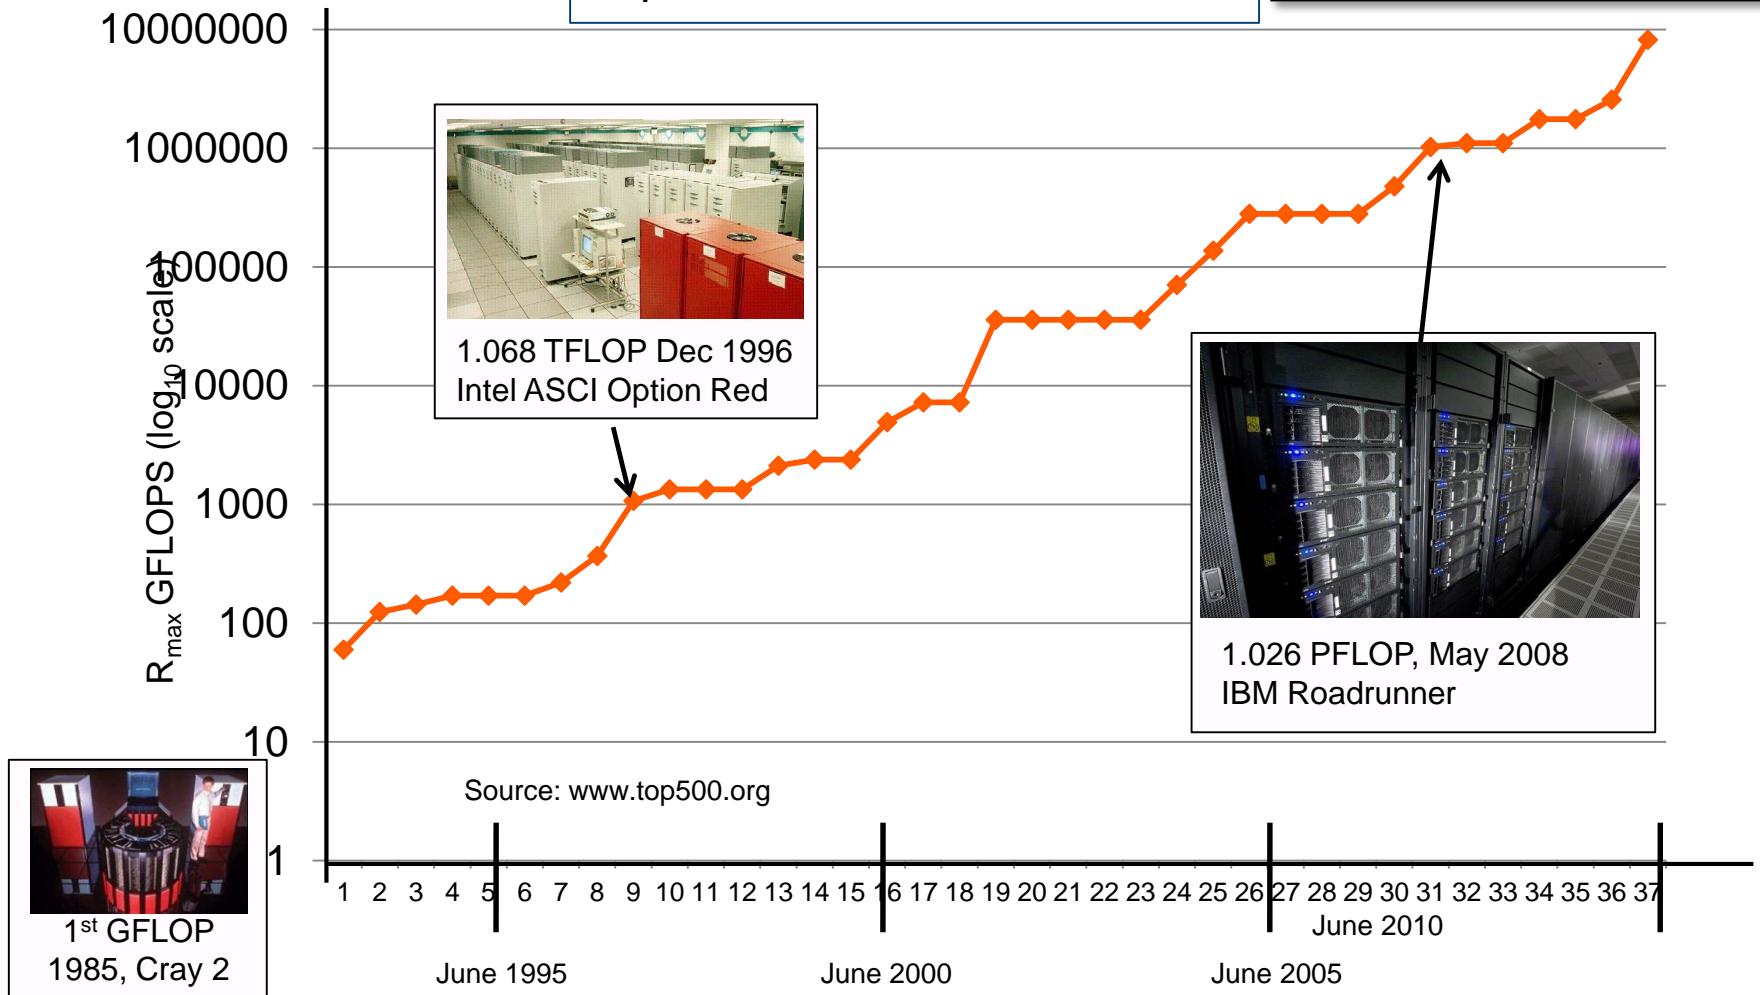
Are there any HPC inflection points looming on the horizon?

Maybe ... consider Exascale computing

# A Linpack ExaFLOP

If we stick to historical trends, we should have that ExaFLOP by 2018

Top500, #1 1993—2011



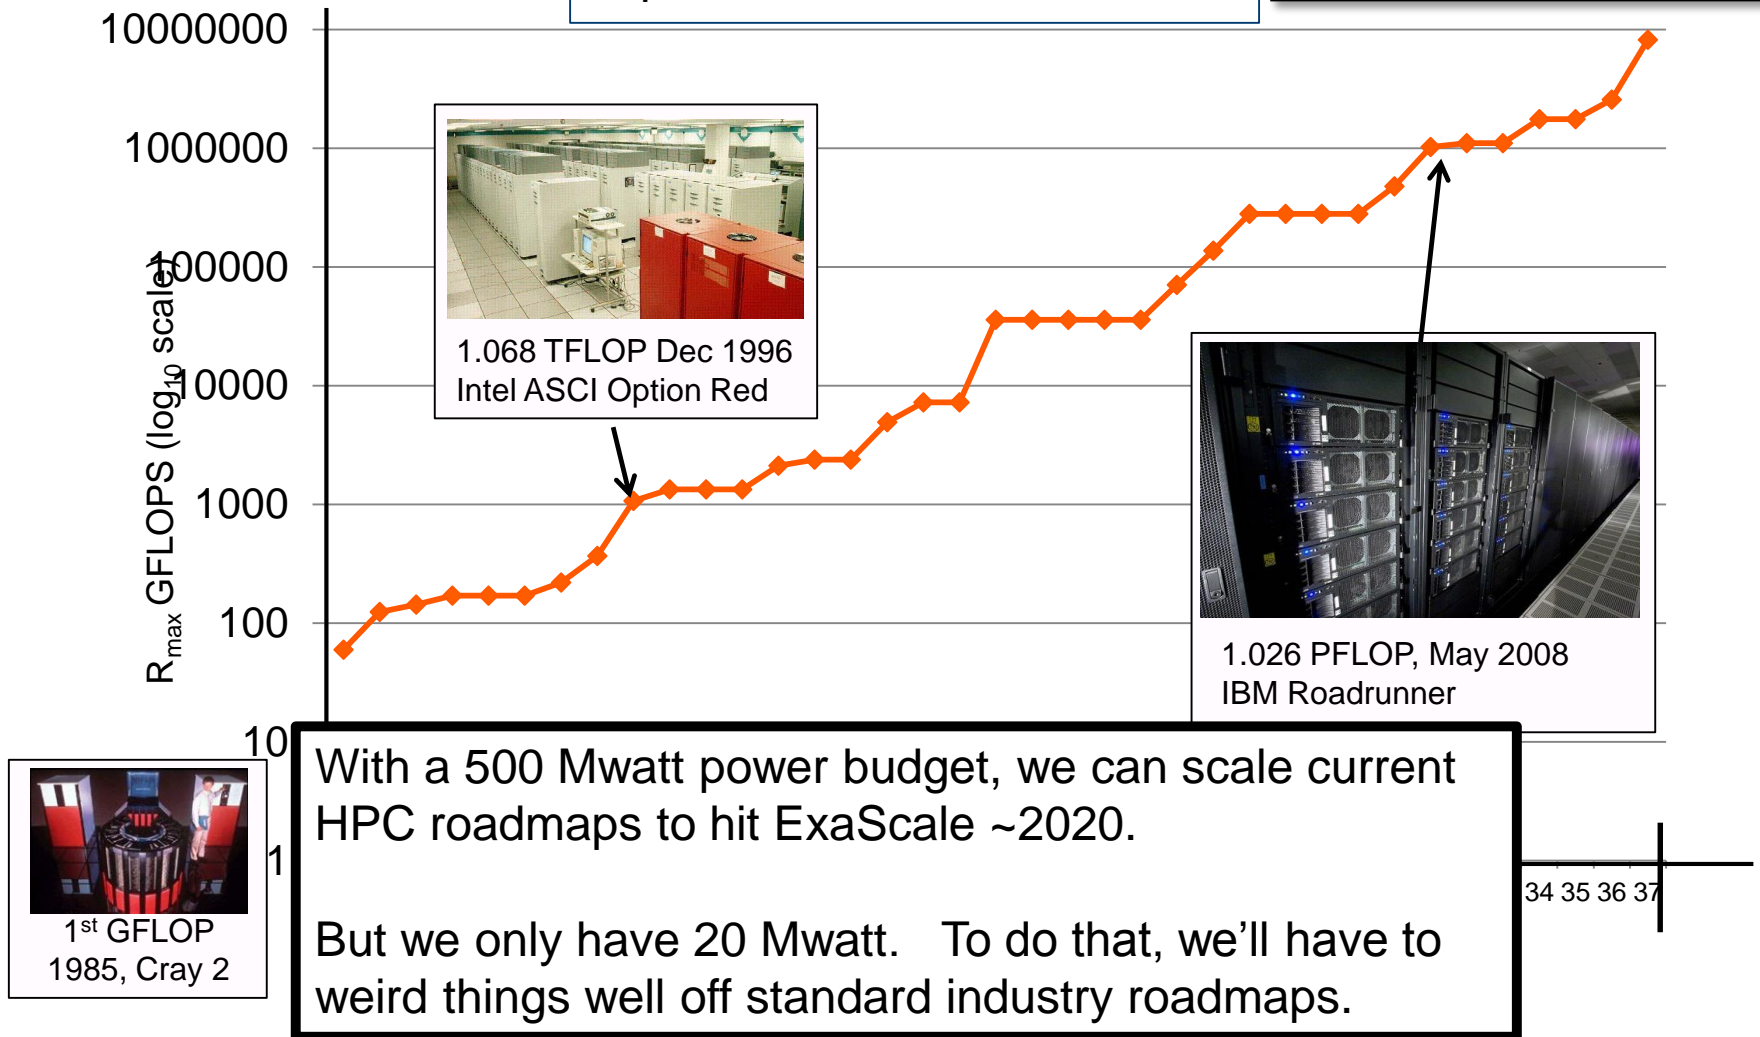
So how are we doing? Are we on track to hit our goal?



# A Linpack ExaFLOP

If we stick to historical trends, we should have that ExaFLOP by 2018

Top500, #1 1993—2011

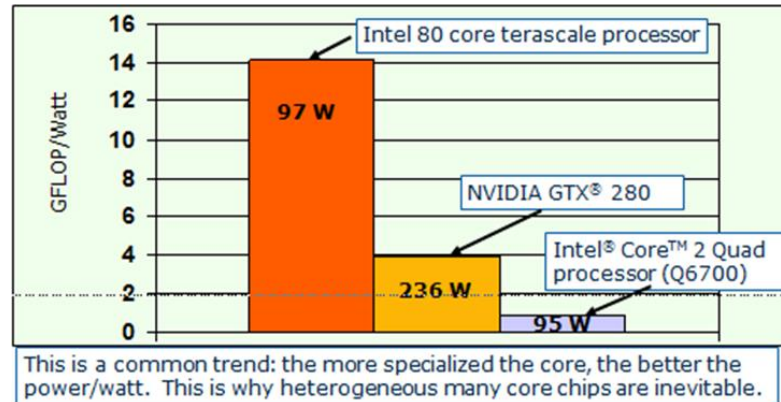


# The ExaScale Challenge: Pity the Poor Programmer

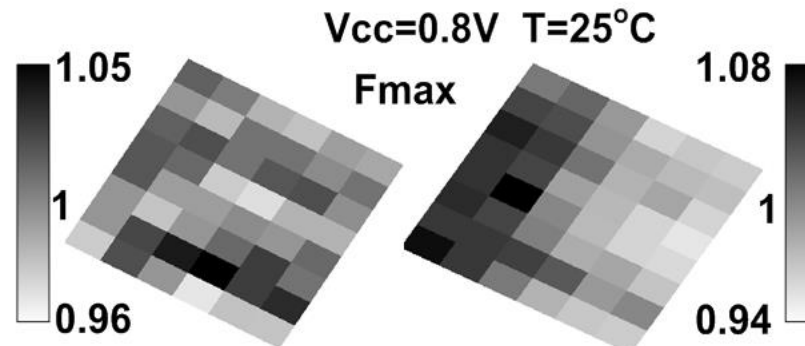
- Vast numbers of low power cores

$$\begin{aligned}\text{Capacitance} &= C \\ \text{Voltage} &= V \\ \text{Frequency} &= f \\ \text{Power} &= CV^2f\end{aligned}$$

- A mix of general purpose, vector, and special function cores.



- Near threshold logic at 7 nm (or lower) process technology



Freq. variation across a single many-core chip

- O(billion) parallelism. Remember Amdahl's law?
- Parallel composition of SW specialized to different cores
- Variability, reliability, and silent errors

# Probability of Faults

Fault	Probability	Impact	Action
Fans	High	Low	Node down
Power Supply	High	Low	Node down
CPU / SRAM	Very Low	Low	Node down
DRAM	Medium	Low	Reconfiguration
Solder Joints	High	Low	Node down
Sockets	High	Low	Node down
Disks	Mid to High	Low	Reconfiguration
NAND/PCM	Low	Low	Reconfiguration
Soft Errors	Low	High	Clever accounting

Source: John Daly, David Mountain's NSA Resiliency WS 02/2012

Source: Shekhar Borkar, Intel, Sept. 2014.

# Probability of Faults

Fault	Probability	Impact	Action
Fans	High	Low	Node down
Power Supply	High	Low	Node down
CPU / SRAM	Very Low	Low	Node down
DRAM	Medium	Low	Reconfiguration
Solder Joints	High	Low	
Sockets	High	Low	
Disks	Mid to High	Low	
NAND/PCM	Low	Low	
Soft Errors	Low	High	accounting

Only a HW person would mark these low ... to my app, these kill the program and therefore have a high impact

accounting

Source: John Daly, David Mountain's NSA Resiliency WS 02/2012

Source: Shekhar Borkar, Intel, Sept. 2014.



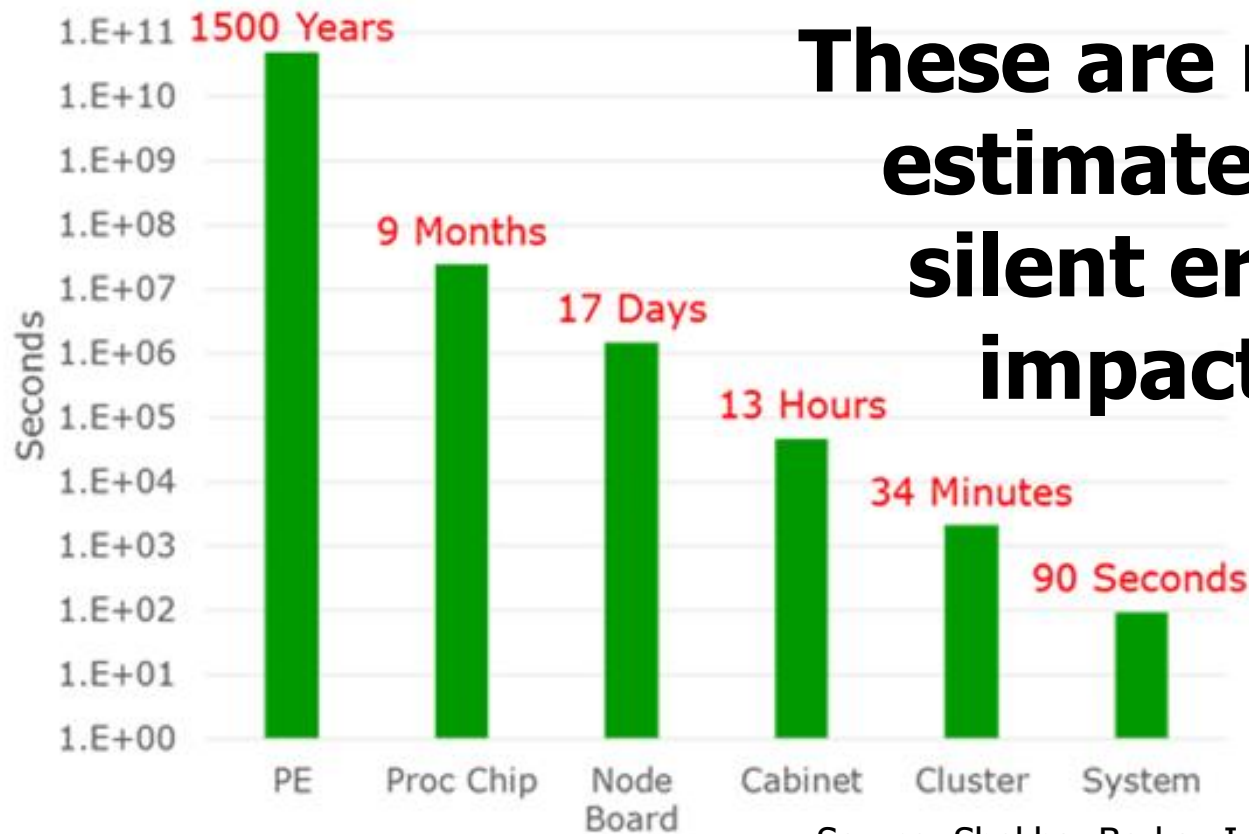
# Probability of Faults

Fault	Probability	Application Software	Action
		Impact	
Fans	High	High	Node down
Power Supply	High	High	Node down
CPU / SRAM	Very Low	High	Node down
DRAM	Medium	High	Reconfiguration
Solder Joints	High	High	Node down
Sockets	High	High	Node down
Disks	Mid to High	High	Reconfiguration
NAND/PCM	Low	High	Reconfiguration
Soft Errors	Low	Terrifying	Clever accounting

Source: John Daly, David Mountain's NSA Resiliency WS 02/2012

Source: Shekhar Borkar, Intel, Sept. 2014.

# Mean Time to Error (Single Event)



Source: Shekhar Borkar, Intel, Sept. 2014.

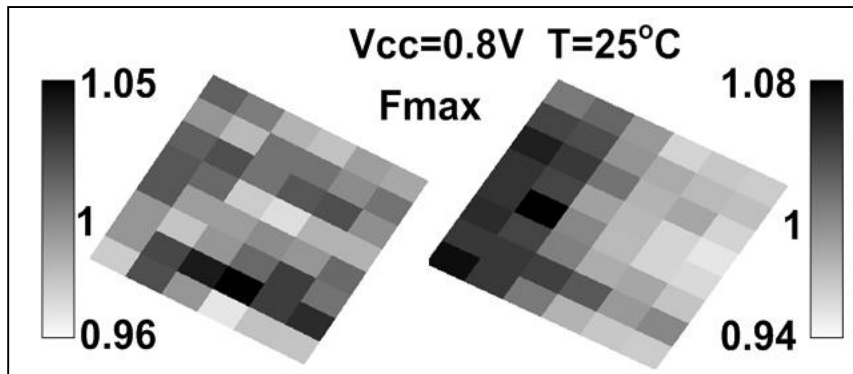
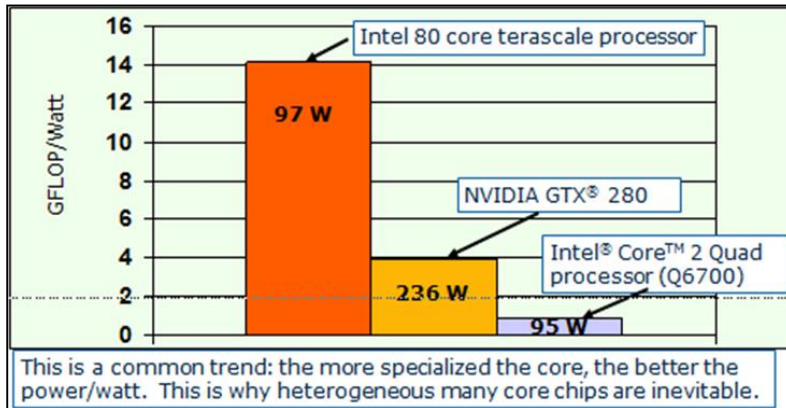
**Select confinement and recovery (state restore) based on probability of error**

i.e. an exascale system is a neutron detector every 90 seconds

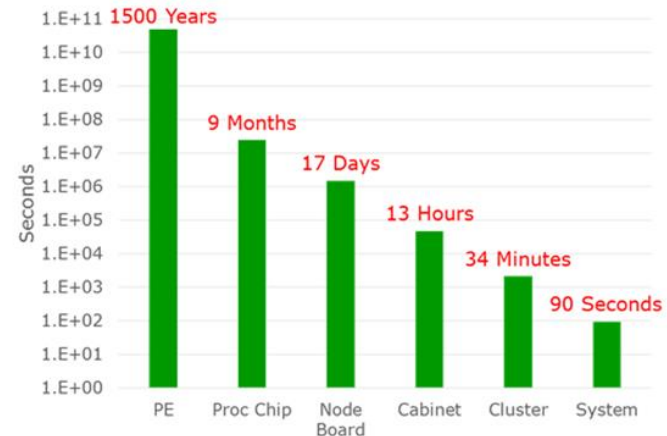
Assumptions,  $10^{-4}$  soft error rate for an exaScale system with 2000 PE/cores per proc, 16 procs per node, 32 nodes per cabinet, 500 Cabinets per system

# HW inflection points

- Resilience + heterogeneity + extreme scale suggests exaScale will be a major hardware inflection point.



## Mean Time to Error (Single Event)



Select confinement and recovery (state restore) based on probability of error

So maybe at last, we can realistically push something besides MPI+X into the HPC applications community?

# 2 pathways to Exascale Runtime Research

Evolutionary  
(e.g. MPI+X)

Revolutionary  
(e.g. OCR)



## Systemic Exascale Challenges

System Utilization

Asynchrony

Data movement cost

Load Imbalance

Fault Tolerance

Scalability

# Exascale Runtime Candidates

Model	Memory	Load balancing		User interface	Execution model	Fault tolerance	Collectives
		SMP	Between nodes				
<b>OCR</b>	relocatable data blocks	work stealing	data block seeding	library	event driven tasks	----	barrier
<b>Charm++</b>	relocatable (arrays of) objects	migration	migration + object seeding	translator	asynch RMI	local checkpoint + message logging	barrier, bcast, reduce, chare array
<b>Legion</b>	relocatable hierarchical data blocks	work stealing	data block seeding + ?	translator or library	asynch RMI	----	barrier, reduce, task array
<b>Grappa</b>	PGAS	work stealing		library	data driven tasks	----	barrier, bcast, reduce, on-all-cores
<b>HPX/ XPI</b>	AGAS/PGAS	work stealing	---	library	message driven execution	----	barrier, reduce
<b>MPI+X</b> (X= MPI OpenMP)	Node-local + X	X	App-level	Library + X	CSP+X	UFLM	Full set

Third party names are the property of their owners

# Exascale Runtime Candidates

Model	Memory	Load balancing		User interface	Execution model	Fault tolerance	Collectives
		SMP	Between nodes				
<b>OCR</b>	relocatable data blocks	work stealing	data block seeding	library	event driven tasks	----	barrier
<b>Charm++</b>	relocatable (arrays of) objects	migration	migration + object seeding	translator	asynch RMI	local checkpoint + message logging	barrier, bcast, reduce, chare array
<b>Legion</b>	Often referred to as Asynchronous Multi-tasking (AMT) systems				asynch RMI	----	barrier, reduce, task array
<b>Grappa</b>	PGAS	work stealing		library	data driven tasks	----	barrier, bcast, reduce, on-all-cores
<b>HPX/ XPI</b>	AGAS/PGAS	work stealing	---	library	message driven execution	----	barrier, reduce
<b>MPL-X</b> (X= MPI OpenMP)	Node-local + X	X	App-level	Library + X	CSP+X	UFLM	Full set

# You can only understand a software system by using it.

- Ideal Case ...
  - This is an important problem, so we must take as long as needed to get it right
  - You must work with full applications
  - Many problems only appear “at scale”
- The harsh byte of reality
  - Hardware is only a few years away and applications teams need to start soon to get ready. We need answers NOW!!!
  - Full applications take Person-YEARS to write/optimize for a new system. With so many systems to study, we simply must find simpler test cases to use.
  - Full Scale systems are busy doing real work.

# The Parallel Research Kernels (PRK)

**PRK: Low level constructs that capture the essence of what parallel programmers require from parallel computers.**

- **The PRK**
  - A set of low level “research kernels” to guide the design of future systems.
  - Output from full apps analysis overwhelming, hard to digest. The PRK can even be understood by HW engineers.
  - They are simple, small and easy to run ... support simulator analysis or even “paper and pencil” machines.
- **Selection Methodology**
  - Anecdotal analysis by a panel of experienced parallel application programmers at Intel.
  - Great Care was taken to make sure their experience covered the full range of typical HPC applications.



# The Parallel Research Kernels version 1.0

- Dense matrix transpose
- Synchronization: global (collective) and point to point
- Scaled vector addition (Stream triad)\*
- Atomic reference counting, both contended and uncontended (locks/TSX)
- Vector reduction
- Sparse matrix-vector multiplication
- Random access update
- Stencil computation
- Dense matrix-matrix multiplication (~DGEMM)
- Branch (inner-loop conditionals + PC panic)\*

These are for “old school” HPC; can be statically load balanced.  
Summer’2016 we’ll add a new case that demand dynamic load  
balancing (a Particle in Cell code)

# The Parallel Research Kernels version 1.0

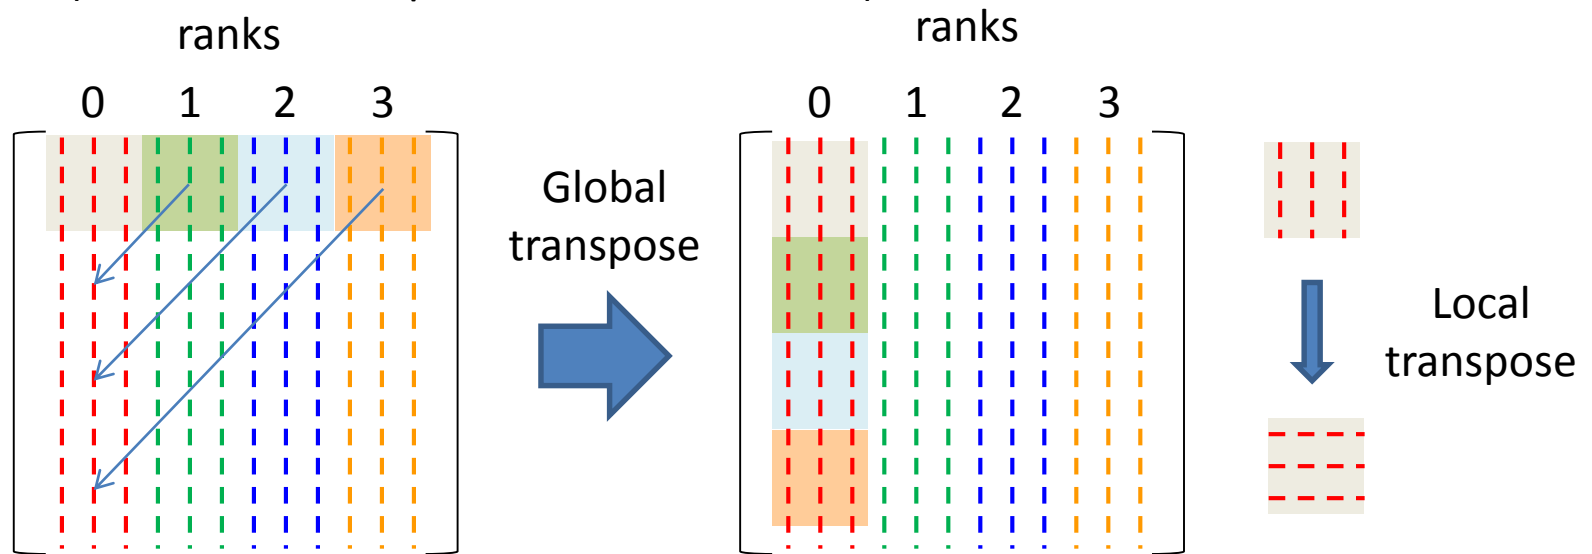
- **Dense matrix transpose**
- **Synchronization**: global (collective) and **point to point**
- Scaled vector addition (Stream triad)\*
- Atomic reference counting, both contended and uncontended (locks/TSX)
- Vector reduction
- Sparse matrix-vector multiplication
- Random access update
- **Stencil computation**
- Dense matrix-matrix multiplication (~DGEMM)
- Branch (inner-loop conditionals + PC panic)\*

**Red Text** indicates  
subset used in our  
exaScale studies

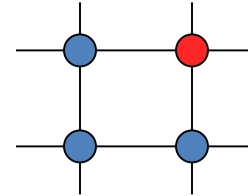
These are for “old school” HPC; can be statically load balanced.  
Summer’2016 we’ll add a new case that demand dynamic load  
balancing (a Particle in Cell code)

# Dense matrix Transpose

- What:  $A = B^T$ , both A and B distributed identically, and by whole columns, using column-major storage format
- Why: Stride 1 on read, large stride on write (TLB, cache misses). All to all communication. Common operation in HPC (FFTs, numerical linear algebra)
- How: Implemented in MPI, MPI+OpenMP, MPI+MPI3-shared memory, Charm++ (tiled local transpose to reduce misses)



# Point to point synchronization (Synch\_p2p)



What:  $A(i,j) = A(i-1,j) + A(i,j-1) - A(i-1,j-1)$

Pipelined solution of problem with non-trivial 2-way dependencies

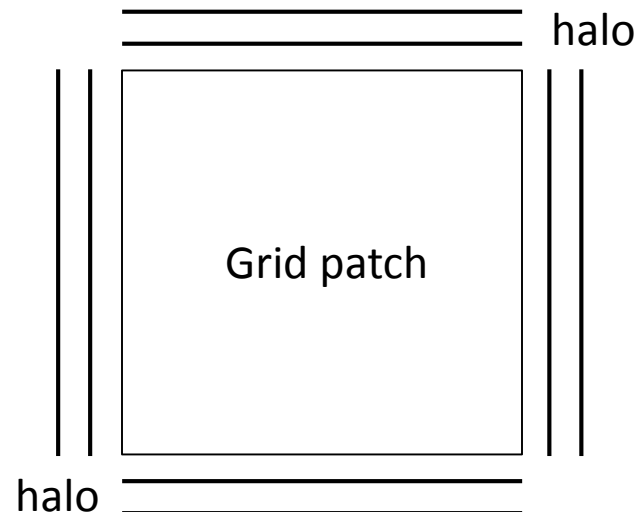
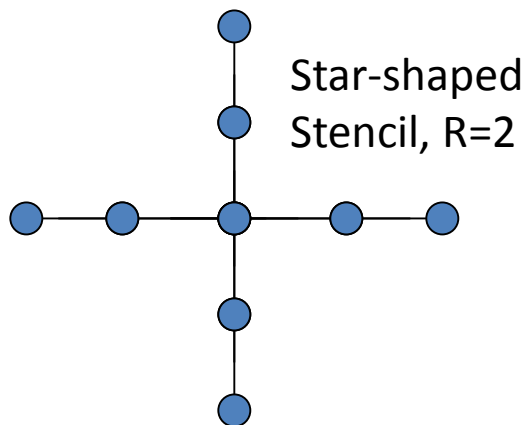
Why: Threads/ranks forced to synchronize pairwise frequently during each iteration. Stresses fine grain communication/synchronization.

How: Implemented in MPI, MPI+OpenMP, MPI+MPI3-shared memory, Charm++, Grappa



# Stencil computation

- What: For all points in 2D grid, compute  $a = S(b)$ ,  $S$  is star-shaped stencil operator (radius  $R$ ),  $a$  and  $b$  are scalar grid variables (2D arrays)
- Why: Data parallel; point-to-point bulk communication; common operation in HPC
- How: Implemented in MPI, OpenMP, MPI+MPI3-SHM, MPI+OpenMP, Charm++, Grappa; use 2D decomposition, collect halo values at start of each iteration



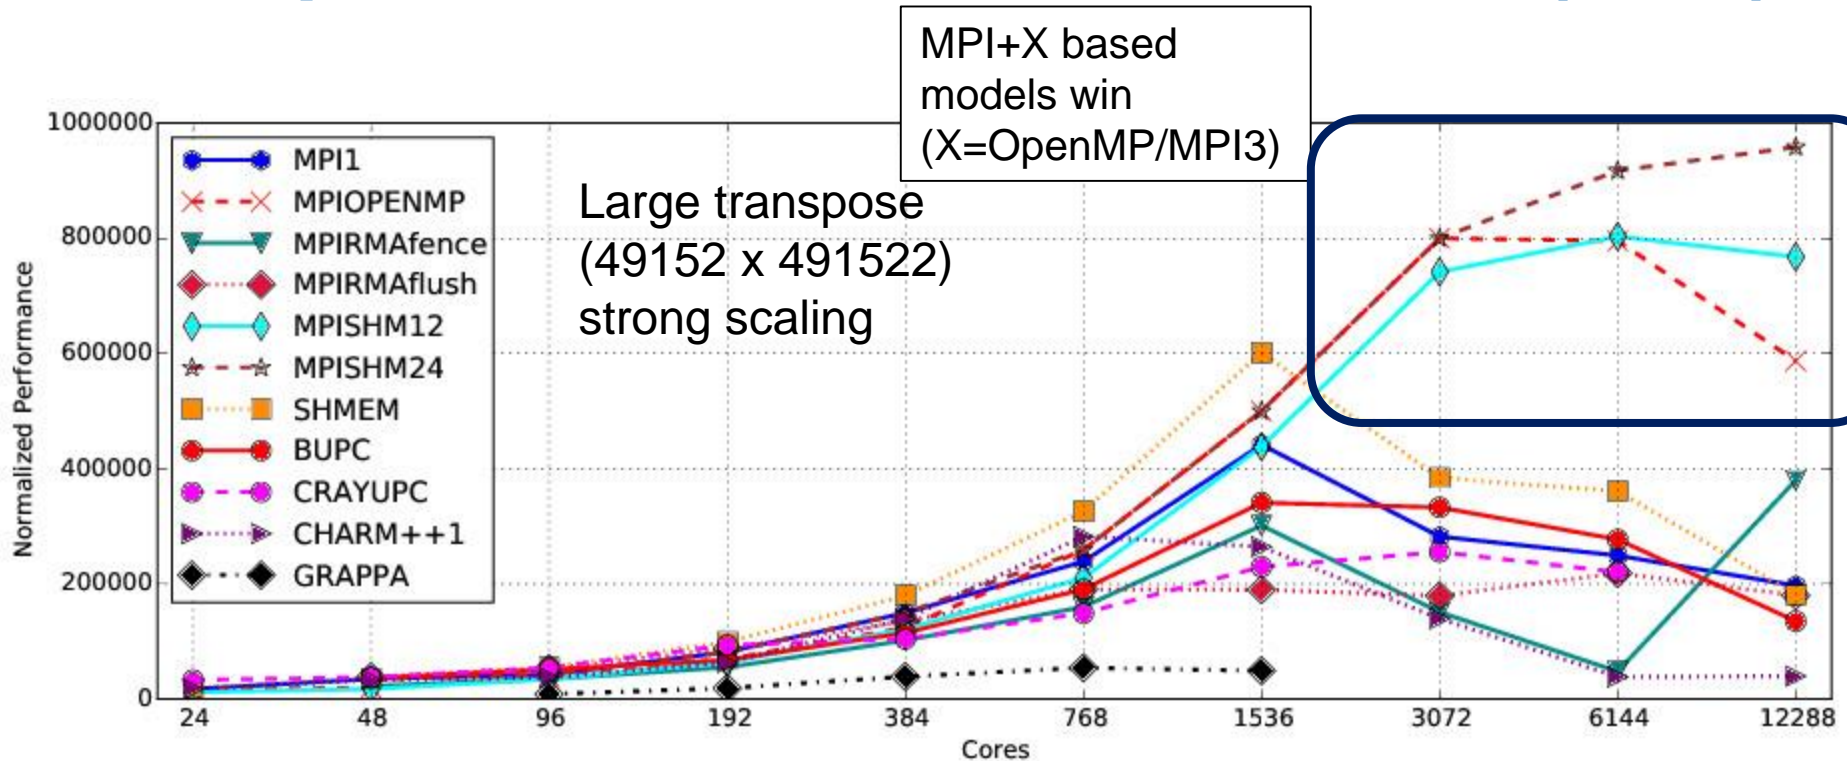


# Experimental apparatus



Cray XC30 Supercomputer with two 12-core Intel® Xeon® E5-2695 processor per node with the Aries interconnect in a Dragonfly topology. Intel compiler version 15.0.1.133 for all codes except Cray compiler environment (CCE 8.4..0.219 for Cray UPC and gcc 4.9.2 for Grappa. Cray MPT 7.2.1 for MPI and SHMEM

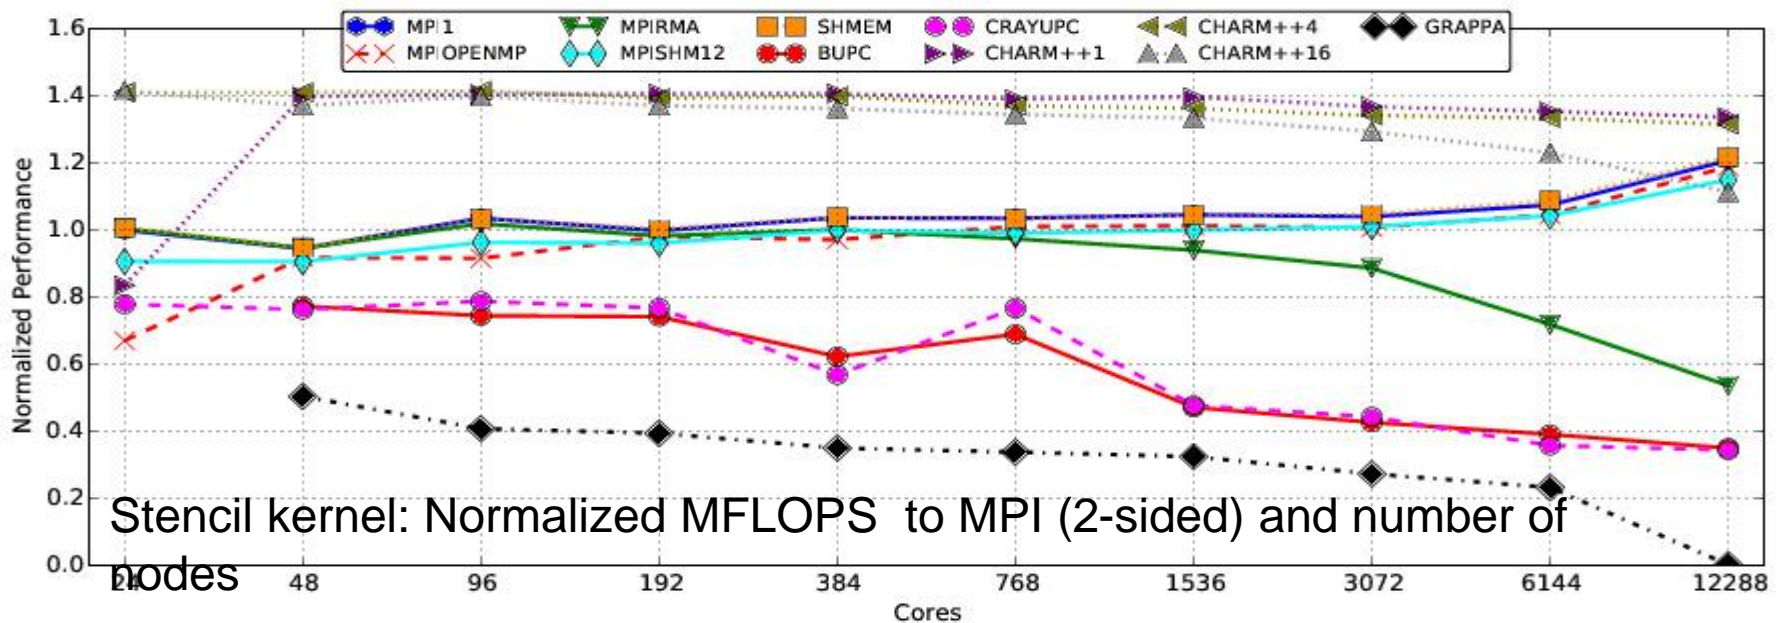
# Transpose Parallel Research Kernel (PRK)



- MPI-OpenMP, MPI-SHMEM do very well since this coarse grained kernel is ideally suited to their SPMD pattern
- Grappa generates large numbers of small messages as the core-count increases .. hence its performance is poor.



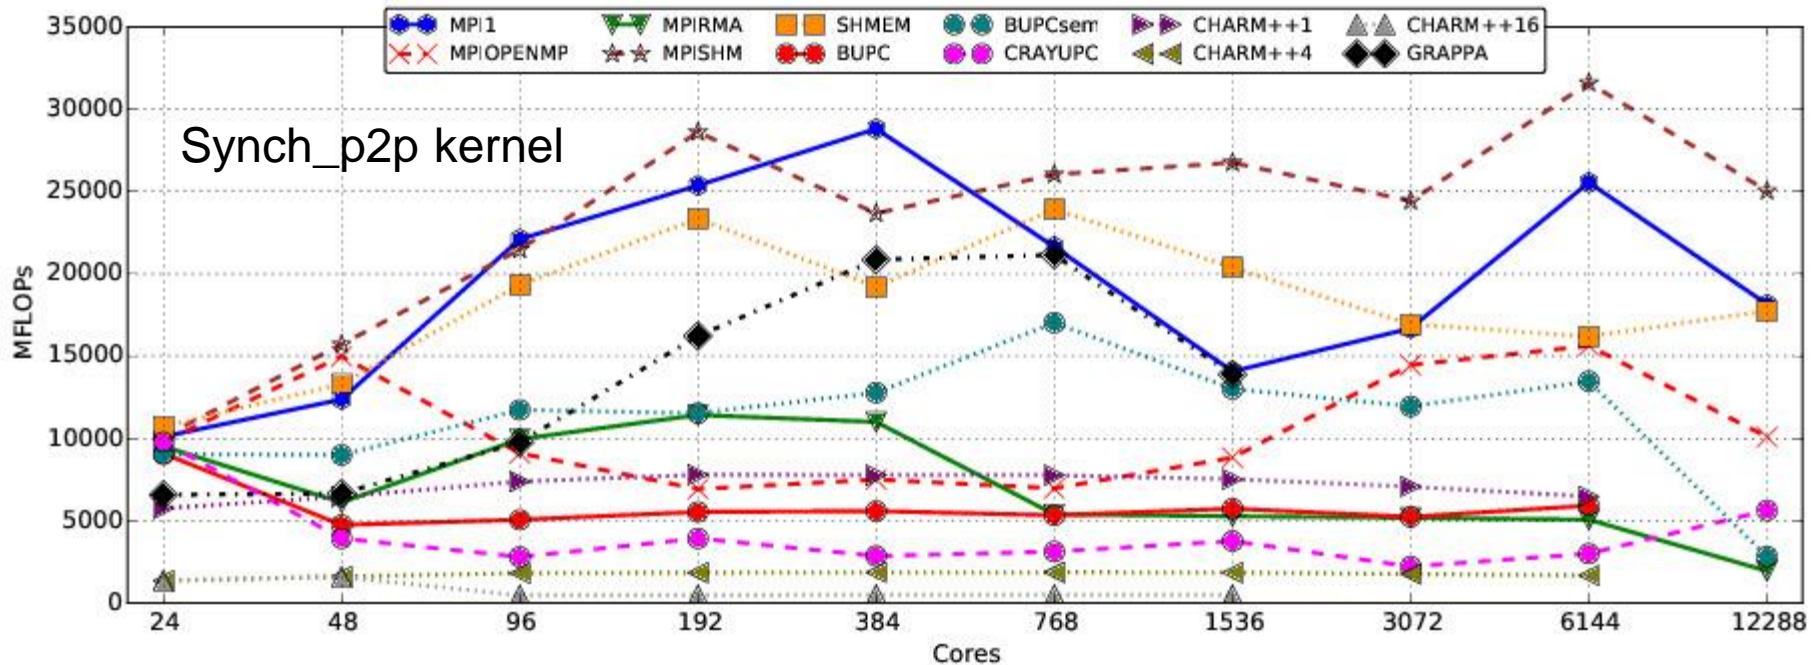
# Stencil Parallel Research Kernel (PRK)



- Runtimes that depend on barriers (MPIRMA, BUPC, CrayUPC) do not do well.
- Runtimes that use pair wise sync (Charm++, Grapa, SHMEM, MPI1) scale well.
- Charm++ handles this example particularly well (numbers 1, 4, and 16 are the degrees of over decomposition)

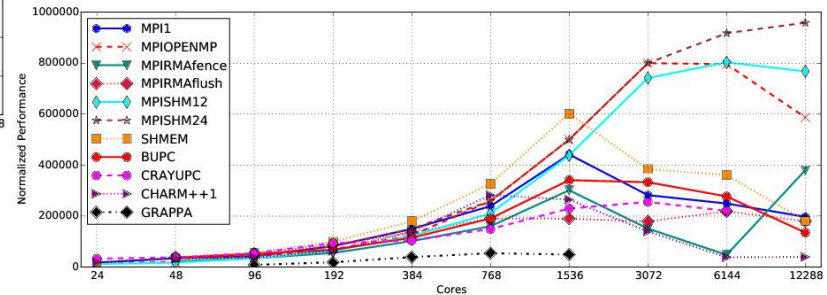
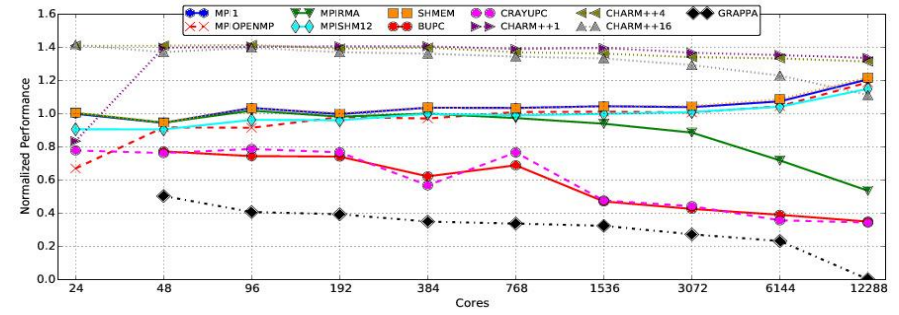
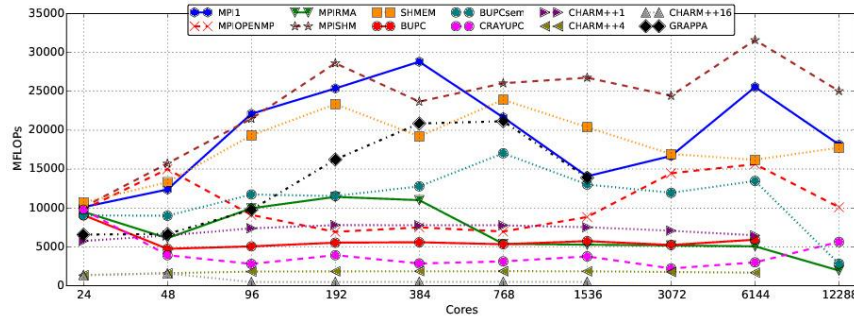


# Synch\_p2p Parallel Research Kernel (PRK)



- Performance dominated by (1) frequent pair-wise synchronization and (2) global sync after each iteration.
- MPIRMA and both versions of UPC used barriers after each iteration ... this really hurt their performance.
- MPI1, SHMEM, and MPISHM have very similar perf.
- This workload is poorly suited to the dynamic execution profile charm++ was designed for.

# How much should we emphasize performance



- Modern Asynchronous Multitasking (AMT) Systems are immature compared to MPI and OpenMP.
  - It's difficult to separate pathological problems with the model from low immature implementations of models.
- What really matters is programmability ... how effective are these programming models for application programmers?

# Programmer Experience notes: Legion

- No language spec, no documentation, no user-manual.
  - Only doxygen based on source code documentation
  - Hard to identify, discover some of the runtime API
- Need to learn too many new terms, structures specific to Legion
  - Logical regions, physical regions, fields, index spaces etc.
- Physical Region accessors are not intuitive, no direct array based indexing, e.g. `value = acc_a[i]`
- Due to the live development, sometimes hard to decide whether you hit a Legion runtime bug or not
- Collectives are not user friendly as advertised
  - IndexSpace tasks should have a direct reduction operation on the calculated Future value but API requires a lot of verbosity to implement this in the code
  - Again finding the correct API, best method is cumbersome
  - Still working on resolving some ambiguities on this

# Programmer Experience notes: Charm++

- Not having a Charm++ compiler greatly hurts productivity.
  - Many programmer errors result in *auto-generated header files* that trigger compiler errors.
  - Interface files (.ci) have special syntax that looks like C++, but isn't.
  - Interface file parser needs to be expanded each time an unexpected construct is encountered—result of no interface file language specification.
  - Collectives in interface file must have different names, even if they do identical things. Consequence: reusing code is harder than it should be.
  - Global variables must be declared in .C file **and** in interface file. Failure to declare these as *readonly* in interface file leads to runtime errors, not compiler errors.
  - Having to declare code as *structured dagger* code is incomprehensible to the programmer. It is just Charm++.
  - Making the programmer responsible for functions **and** their proxies (must maintain two sets of names) is awkward at best, and invites errors.
- Building Charm++ on top of MPI and/or using non-gnu tools not obvious.

# Programmer Experience notes: Grappa

- No user manual.
- No “threadprivate” variables, except through global variables.
- Support for global variables weak/buggy.
- Only two types of synchronization: barrier, and full-empty bits; leads to inefficient code.
- Full-empty bits implementation incomplete.
- Compile and link flags not statically determined; must be copied from special location after Grappa build.
- Building Grappa with Cmake not obvious.
- Having to support external repos or hack around them cumbersome.

# Programmer Experience notes: OCR

- The specification is GREAT!!!
- For something as old as OCR, the implementation on scalable systems is way behind the competition.
- It is not programmer friendly ... yes it's a runtime and not designed for application programmers, but still, it shouldn't make "the easy things so hard"

# Advice to the AMT community

- AMT is like communism ...
  - Its a great idea but we just don't have any good implementations to reference.
- Doxygen is NOT useful documentation for applications programmers.
  - Write a specification ... It will be good for you
  - Write a programmers reference manual
  - Create an official examples suite to help people learn.
- Don't fight MPI on it's own turf. The advantage of AMT that is hard for MPI+OpenMP to match is resilience. So why is it that only Charm++ has a mature resilience implementation?
- Remember if you compete on MPI's turf, (paraphrasing the old joke about hikers and bears ..)
  - *MPI does not need to outrun the bear (exascale requirements), just the other hikers (disruptive runtimes), to be a winner*



# Ease of use

- Once you climb the MPI learning curve, it's a natural and easy model to work with:
  - The people I hear complaining the most about MPI are the ones who don't use it ... often AMT researchers.
  - MPI programmers like MPI!!! It runs everywhere and is easy to debug.
- Key change I'd like to see in the AMT community:
  - Stop making the things that are easy in MPI/OpenMP so hard:
    - Collectives
    - Launching large SPMD programs across a scalable system
  - Legacy code is far more important than many people think. If your solution to my programmability program begins with the phrase .. "Recode everything in my new language", I'll stop listening to you.
- A good AMT model must make it easy for me to program "in the large" with algorithms I care about, not the ones that nicely match what your programming model expects.

# 2 pathways to Exascale Runtime Research

Evolutionary  
(e.g. MPI+X)

Revolutionary  
(e.g. OCR)

Based on our experiences with AMT systems and understanding of you application programmers work, we are choosing a path to exascale based on MPI.

## Systemic Exascale Challenges

System Utilization

Asynchrony

Data movement cost

Load Imbalance

Fault Tolerance

Scalability

# MPI for Darwinists:

An extended MPI that coherently integrates disruptive new features

## 1. Fault Tolerance in MPI

- Local checkpoint/restart for task rollback, ability to serialize and migrate ranks. (collaborators: Fenix and UFLM)

## 2. Transparent Over decomposition in MPI:

- Ranks-as-threads integrated into the MPI runtime. (collaborators: Fine Grained MPI Group)

## 3. Infrastructure to support load balancing

- Rank migration (collaborator: AMPI)

## 4. Investigate new, light-weight load balancing schemes (on and between nodes):

- MPI+MPI: Rank stealing
- MPI+X: OpenMP 4+ (tasks, target), TBB

# But this is all just anecdotal evidence. Is that really enough?

- Programmability is too important to leave to anecdote.
- We need a systematic framework to access programmability.
  - HPCS program made a noble attempt to create such a framework:
    - Measured actual programmer productivity
    - Found data showing that (1) once a project chooses a language, they won't change; (2) OpenMP is NOT harder to debug than MPI; (3) programmers care about portability, performance and the ability to make incremental changes more than elegance.
    - Constructed formal models including some based on economic utility theory.
  - The HPCS program, however, had no lasting impact on programmability. Why? My untested hypothesis
    - They didn't engage cognitive psychologists in their studies of programmers and hence didn't not come up with models based on how human's think about programming ... it was hard to turn the HPCS results into action

# Psychology of Programming in one slide

- Human reasoning is model based ... Programming is a process of successive refinement of a problem over a hierarchy of models.[1]
- Programmers use an informal, internal notation based on the problem, mathematics, programmer experience, etc.
  - Within a class of programming languages, the solution generated is only weakly dependent on the language.[2] [3]

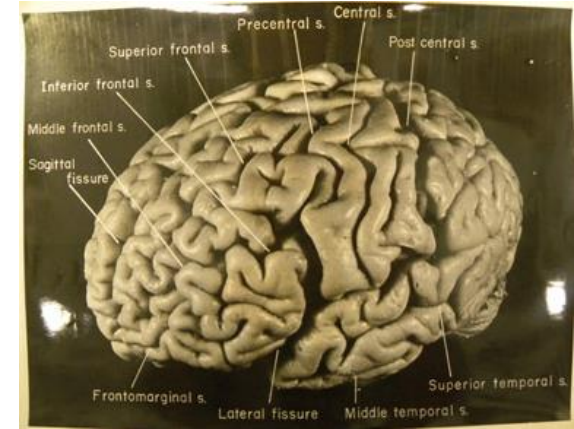


Image source: Einstein's brain from <https://scim.ag/Einsteinbrain>. Collected on 12/4/2012

- Programmers think about code in chunks or “plans”. [4]
- Opportunistic Refinement: [5]
  - Progress is made at multiple levels of abstraction with effort focused on the most productive level. It's not top-down or bottom-up .... We bounce!

[1] R. Brooks, "Towards a theory of the comprehension of computer programs", Int. J. of Man-Machine Studies, vol. 18, pp. 543-554, 1983.

[2] S. P. Robertson and C Yu, "Common cognitive representations of program code across tasks and languages", int. J. Man-machine Studies, vol. 33, pp. 343-360, 1990.

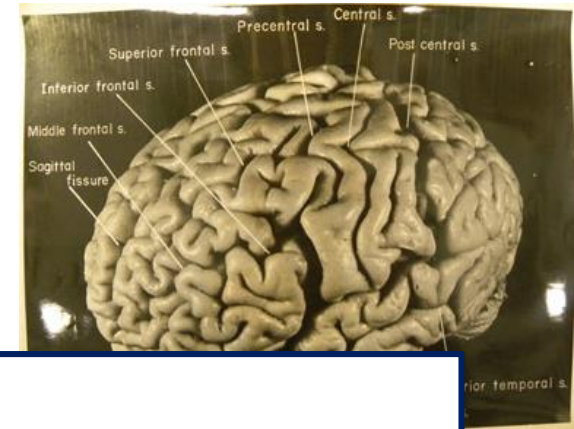
[3] M. Petre and R.L. Winder, "Issues governing the suitability of programming languages for programming tasks. "People and Computers IV: Proceedings of HCI-88, Cambridge University Press, 1988.

[4] R.S. Rist, "Plans in programming: definition, demonstration and development" in E. Soloway and S. Iyengar (Eds.), Empirical Studies of Programmers, Norwood, NJ, Ablex, 1986.

[5] M. Petre, "Expert Programmers and Programming Languages", in [Hoc90], p. 103, 1990.

# Psychology of Programming in one slide

- Human reasoning is model based ... Programming is a process of successive refinement of a problem over a hierarchy of models.[1]
- Programmers use an informal, internal notation based on the problem, mathematics, programmer experience, etc.
  - Within a class of programming languages, the



This implies that we want programming models that:

- Make each layer in the hierarchy of models explicitly visible.
- Abstract lower layers but don't hide them.
- Programmers think in "plans" so make them explicitly visible.
- Programmers must be able to control the level of abstraction they work within ... and they need to bounce

[1] R. Brooks, "T

[2] S. P. Roberts  
343-360, 1990.

[3] M. Petre and R.L. Winder, "Issues governing the suitability of programming languages for programming tasks. "People and Computers IV: Proceedings of HCI-88, Cambridge University Press, 1988.

[4] R.S. Rist, "Plans in programming: definition, demonstration and development" in E. Soloway and S. Iyengar (Eds.), Empirical Studies of Programmers, Norwood, NJ, Ablex, 1986.

[5] M. Petre, "Expert Programmers and Programming Languages", in [Hoc90], p. 103, 1990.

# Conclusions

- Programming model developers
  - Innovate and publish, but your output should be to improve/evolve existing standards.
  - Exception ... watch for hardware inflection points. That's the only time new models flourish
  - Sociology is as important as technology ... choice overload and “worse is best” are real.
  - We need to develop a formalism for programmability grounded in the psychology of programming.
- Some key Results with the parallel research kernels
  - Flat MPI often performs best
  - MPISHM usually close, sometimes much better (Transpose); gives greatest flexibility in tuning granularity
  - MPI+OpenMP in canonical form never competitive
  - Charm++ competitive only for medium grain workloads; little/no benefit seen of over-decomposition
  - Legion ... we've been working with Legion for almost two years with little to show. Brutal learning curve of a moving target